

Machine Learning Model

Muhammad Aryaan Lakhani

2025-06-24

```
pacman::p_load( AER, fixest, modelsummary, kableExtra, tidyverse, rsample, # new package for sample splitting  
glmnet, # new package for lasso broom, # new package for model extraction tidymodels # tools for ML in  
tidyverse syntax  
)
```

```
library(AER)
```

```
## Loading required package: car
```

```
## Loading required package: carData
```

```
## Loading required package: lmtest
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##   as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
```

```
## Loading required package: survival
```

```
library(fixest)  
library(modelsummary)  
library(kableExtra)  
library(tidyverse)
```

```
## — Attaching packages  
## —————  
## tidyverse 1.3.2 —
```

```
## ✓ ggplot2 3.4.0      ✓ purrr 1.0.1
## ✓ tibble 3.2.1      ✓ dplyr 1.1.0
## ✓ tidyr 1.2.1       ✓ stringr 1.5.0
## ✓ readr 2.1.3       ✓ forcats 0.5.2
## — Conflicts ————— tidyverse_conflicts() —
## * dplyr::filter() masks stats::filter()
## * dplyr::group_rows() masks kableExtra::group_rows()
## * dplyr::lag() masks stats::lag()
## * dplyr::recode() masks car::recode()
## * purrr::some() masks car::some()
```

```
library(rsample)
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
##
## Loaded glmnet 4.1-7
```

```
library(broom)
library(tidymodels)
```

```
## — Attaching packages ————— tidymodels 1.0.0 —
## ✓ dials 1.1.0      ✓ tune 1.0.1
## ✓ infer 1.0.4      ✓ workflows 1.1.3
## ✓ modeldata 1.1.0  ✓ workflowsets 1.0.0
## ✓ parsnip 1.0.4    ✓ yardstick 1.1.0
## ✓ recipes 1.0.5
## — Conflicts ————— tidymodels_conflicts() —
## * scales::discard() masks purrr::discard()
## * Matrix::expand() masks tidyr::expand()
## * dplyr::filter() masks stats::filter()
## * recipes::fixed() masks stringr::fixed()
## * dplyr::group_rows() masks kableExtra::group_rows()
## * dplyr::lag() masks stats::lag()
## * Matrix::pack() masks tidyr::pack()
## * dplyr::recode() masks car::recode()
## * purrr::some() masks car::some()
## * yardstick::spec() masks readr::spec()
## * recipes::step() masks stats::step()
## * Matrix::unpack() masks tidyr::unpack()
## * recipes::update() masks Matrix::update(), stats::update()
## • Search for functions across packages at https://www.tidymodels.org/find/
```

```
df <- read_csv("/Users/aryaanlakhani/Desktop//Hitters.csv") %>%

# transforming characters to factors
mutate(across(where(is.character), ~factor(.x)))
```

```
## Rows: 263 Columns: 20
## — Column specification —————
## Delimiter: ","
## chr (3): League, Division, NewLeague
## dbl (17): AtBat, Hits, HmRun, Runs, RBI, Walks, Years, CAtBat, CHits, CHmRun...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Data set splitting

```
# setting random seed for replication
set.seed(1999)

# splitting data into training/test
df_split<- initial_split(df, prop = 3/4)

# Creating data
test <- testing(df_split)
train <- training(df_split)
```

Models and Errors

```
models <- list(
  Basic    = lm(Salary ~ Years + RBI + HmRun + CRBI + CHmRun, data = train),
  Full     = lm(Salary ~ ., data = train),
  Extended = lm(Salary ~ (. - Division - League - NewLeague)^2 + Division + League + NewLeague, data = train)
)
```

```
coef(models$Basic) %>% length()
```

```
## [1] 6
```

```
coef(models$Full) %>% length()
```

```
## [1] 20
```

```
coef(models$Extended) %>% length()
```

```
## [1] 140
```

```
models$Basic$rank
```

```
## [1] 6
```

```
models$Full$rank
```

```
## [1] 20
```

```
models$Extended$rank
```

```
## [1] 140
```

```
# Calculating training errors for all models.
```

```
train.mse.Basic <- mean((train$Salary - predict(models$Basic))^2)
```

```
train.mse.Full <- mean((train$Salary - predict(models$Full))^2)
```

```
train.mse.Extended <- mean((train$Salary - predict(models$Extended))^2)
```

```
# Calculating test errors for all models.
```

```
test.mse.Basic <- mean((test$Salary - predict(models$Basic, newdata = test))^2)
```

```
test.mse.Full <- mean((test$Salary - predict(models$Full, newdata = test))^2)
```

```
test.mse.Extended <- mean((test$Salary - predict(models$Extended, newdata = test))^2)
```

```
mse.train <- tibble(
```

```
  msetr = c(train.mse.Basic, train.mse.Full, train.mse.Extended),
```

```
  ktr = c(models$Basic$rank, models$Full$rank, models$Extended$rank)
```

```
)
```

```
mse.test <- tibble(
```

```
  msete = c(test.mse.Basic, test.mse.Full, test.mse.Extended),
```

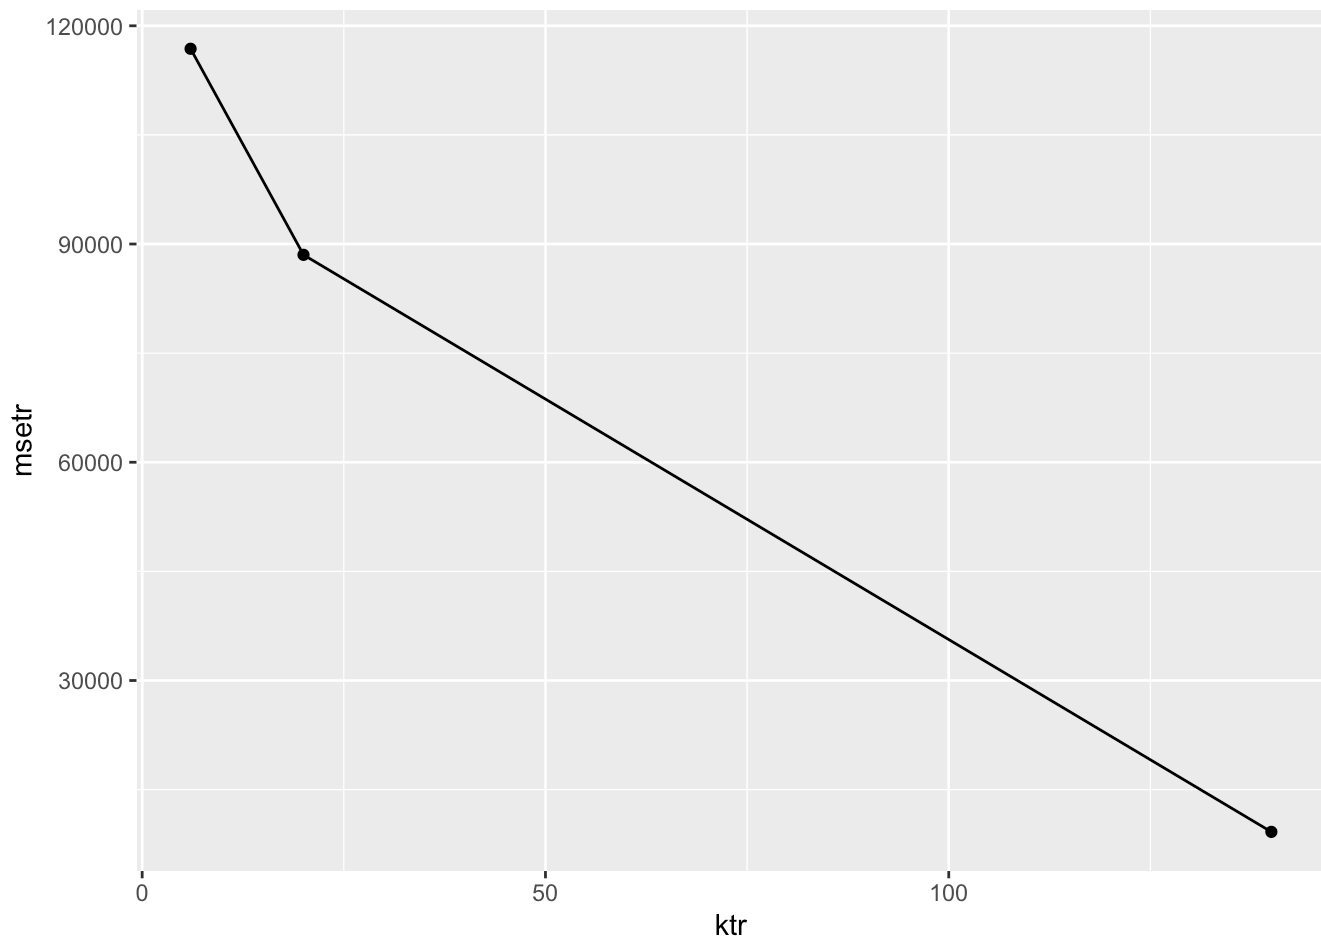
```
  kte = c(models$Basic$rank, models$Full$rank, models$Extended$rank)
```

```
)
```


```
ggplot(mse.train, aes(y = msetr , x = ktr)) +
```

```
  geom_line() +
```

```
  geom_point()
```



```
ggplot(mse.test, aes(y = mse , x = ktr)) +  
  geom_line() +  
  geom_point()
```



Looking at the training dataset, the Basic Model with 6 parameters has the highest MSE, indicating the least accurate fit. The Full Model shows some improvement, with a lower MSE than the Basic Model, but it still doesn't outperform the Extended Model. The Extended Model, with 140 parameters, achieves the lowest MSE, suggesting it fits the training data best and performs the strongest within this dataset.

Examining the models on the test dataset, we observe that the Basic Model performs better here than it did on the training set, as indicated by its lower MSE. In contrast, the Extended Model, which had the lowest MSE during training, shows the highest MSE on the test data—suggesting it may have overfit the training set. The best-performing model on the test data is the Full Model with 20 parameters, achieving the lowest MSE among all the models.

Demonstration

Train data

Extracting Salary as a vector

```
Y_train <- train$Salary
```

Extracting X as matrix

```
X_train <- model.matrix(Salary ~ ., data = train)[, -1]
```

Testing data

Extracting Salary as a vector

```
Y_test <- test$Salary
```

Extracting X as matrix

```
X_test <- model.matrix(Salary ~ ., data = test)[, -1]
```

Ridge

```
ridge <- glmnet(x = X_train,  
               y = Y_train,  
               alpha=0)
```

```
ridge.data <- tidy(ridge)
```

this code chunk is complete, setting eval = T

```
ridge.data %>%
```

```
  filter(term != "(Intercept)") %>%
```

```
  group_by(step) %>%
```

```
  summarize(sum.coef = sqrt(sum(estimate^2)) ,
```

```
            lambda = mean(lambda)) %>%
```

```
  ggplot(aes(y = sum.coef, x = log(lambda))) +
```

```
  geom_line() +
```

```
  labs(
```


```
    y = "l-2 norm",
```

```
    x = "log penalty",
```

```
    title = "Ridge regression: sum of squared coefficients vs penalty"
```


```
)
```

Ridge regression: sum of squared coefficients vs penalty



```
ridge.data %>%  
  filter(term != "(Intercept)") %>%  
ggplot(aes(lambda, estimate, color = term)) +  
  geom_line() +  
  scale_x_log10() +  
  labs(  
    title = "Ridge coefficient path",  
    y = "betas",  
    x = "log lambda"  
  ) +  
  theme(legend.position = 'bottom')
```


Ridge coefficient path

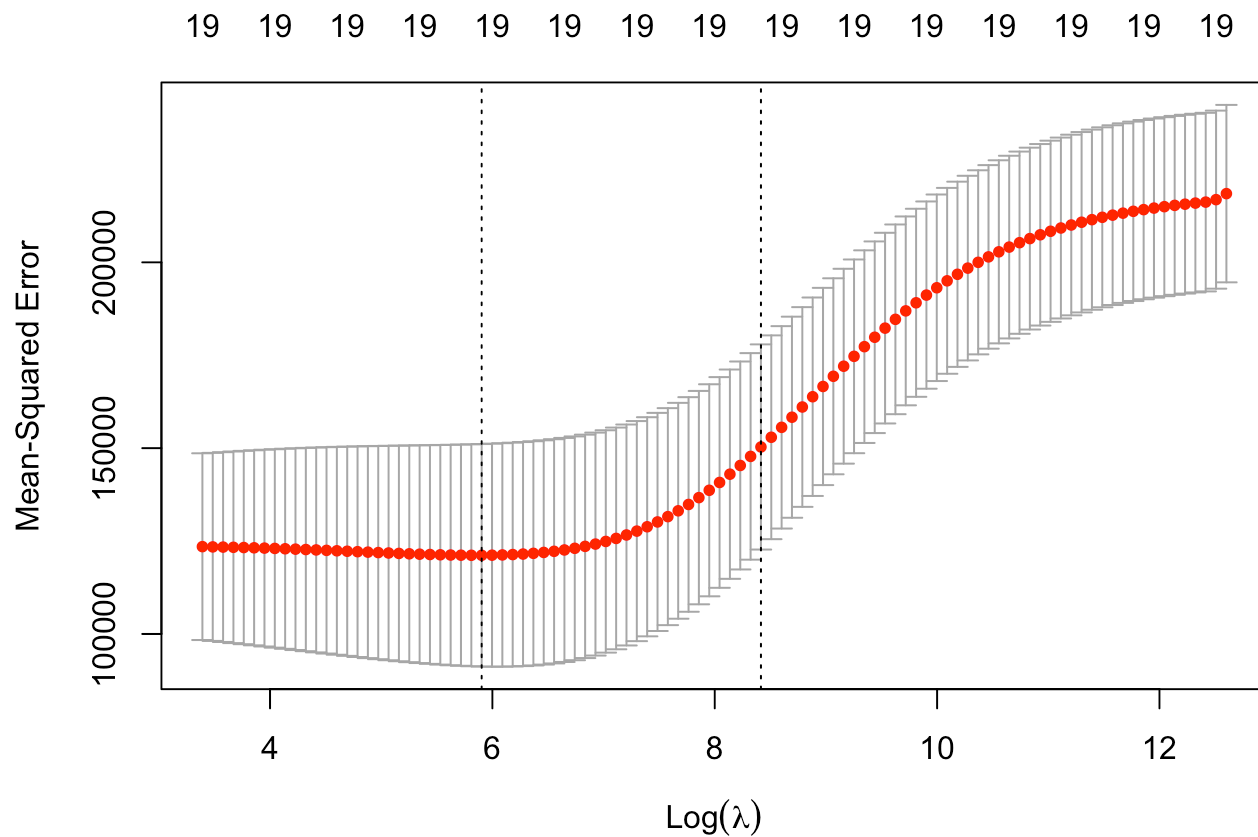


```
cv.ridge <- cv.glmnet(
  y = Y_train,
  x = X_train,
  nfolds = 5,
  alpha = 0)

# output of optimal penalty
cv.ridge
```

```
##
## Call:  cv.glmnet(x = X_train, y = Y_train, nfolds = 5, alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      366    73 121172 29937         19
## 1se     4517    46 150316 27604         19
```

```
# plot of mse vs penalty
plot(cv.ridge)
```



```
# this code chunk is complete, set eval = T
cv.ridge$lambda.min
```

```
## [1] 366.3678
```

```
## [1] 30.7
```

```
# Note the syntax here.
ridge.mse <- mean((Y_test - predict(cv.ridge, s = cv.ridge$lambda.min, newx = X_test))^
2)
ridge.mse
```

```
## [1] 116412.2
```

```
## [1] 129888
```

When comparing the mean squared error (MSE) of Ridge regression to that of the other linear models on the test data, Ridge regression yields a lower MSE—indicating better generalization—i.e., $MSE_{\text{Ridge}} < MSE_{\text{test}}$ for the other models. However, on the training data, the Extended Model achieves the lowest MSE, outperforming Ridge in that context. It's important to note, though, that these models may be optimizing different objectives or using different assumptions, so direct comparisons based solely on MSE may not fully capture their relative performance.

Coding

```
Y_train <- train$Salary
```

```
X_train <- model.matrix(Salary ~ (. - Division - League - NewLeague)^2 + Division + League + NewLeague, data = train)[, -1]
```

```
### Test data
```

```
Y_test <- test$Salary
```

```
X_test <- model.matrix(Salary ~ (. - Division - League - NewLeague)^2 + Division + League + NewLeague, data = test)[, -1]
```

```
cv.ridge <- cv.glmnet(
  y = Y_train,
  x = X_train,
  nfolds = 5,
  alpha = 0)
```


```
ridge.mse <- mean((Y_test - predict(cv.ridge, s = cv.ridge$lambda.min, newx = X_test))^2) %>% mean
ridge.mse
```

```
## [1] 86228.38
```

```
cv.ridge$lambda.min
```


```
## [1] 316.0422
```

```
plot(cv.ridge)
```




```
# lasso
lasso <- glmnet(y = Y_train,
  x = X_train,
  nfolds = 5,
  alpha = 1)

plot(lasso)
```



```
# for the plot
lasso %>%
  tidy() %>%
  filter(term != "(Intercept)") %>%
ggplot(aes(lambda, estimate, color = term)) +
  geom_line() +
  scale_x_log10() +
  labs(
    title = "Lasso coefficient path",
    y = "betas",
    x = "log lambda"
  ) +
  theme(legend.position = "none")
```




#The coefficient path for Lasso regression appears more "spiky" compared to that of Ridge regression. This is due to the nature of the Lasso penalty, which is more aggressive in driving some coefficients to exactly zero. As the regularization parameter (λ) increases, many coefficients in Lasso are completely eliminated, leaving only a few with non-zero values—hence the jagged, abrupt changes in the path. On the other hand, Ridge regression applies a softer penalty that continuously shrinks all coefficients toward zero without eliminating them entirely, resulting in a much smoother and more gradual coefficient path.

```
cv.lasso <- cv.glmnet(
  y = Y_train,
  x = X_train,
  nfolds = 5,
  alpha = 1)

cv.lasso$lambda.min
```

```
## [1] 3.896315
```

```
plot(cv.lasso)
```



```
lasso.mse <- mean((Y_test - predict(cv.lasso, s = cv.lasso$lambda.min, newx = X_test))^2) %>% mean
lasso.mse
```

```
## [1] 76997.48
```

```
ols <- glmnet(
  y = Y_train,
  x = X_train,
  lmabda = 0
)

ols.mse <- (Y_test - predict(ols, s = 0, newx = X_test))^2 %>% mean

results <- tibble(
  model = c("ols", "lasso", "ridge"),
  mse = c(ols.mse, lasso.mse, ridge.mse)
)

results %>%
  kbl() %>%
  kable_classic_2(full_width = F)
```

model	mse
ols	147384.35
lasso	76997.48
ridge	86228.38

```
# print the results
print(results)
```

```
## # A tibble: 3 × 2
##   model      mse
##   <chr>    <dbl>
## 1 ols    147384.
## 2 lasso   76997.
## 3 ridge   86228.
```

Among the models evaluated, the Ridge regression model achieved the lowest mean square error (MSE), indicating the best performance. The Lasso model followed closely behind, while the Ordinary Least Squares (OLS) model had the highest MSE, making it the least effective in this case. Therefore, Ridge regression outperforms the others in this scenario.

The final lasso model includes only a subset of the original variables. The coefficients for some of the variables are exactly zero, indicating that they have been entirely eliminated from the model. The variables that are still present in the model are considered important by the lasso regression algorithm. In the output provided, the intercept is the first coefficient and is equal to 340. The second coefficient corresponds to the variable "AtBat" and has a value of -0.595. The third coefficient is for the variable "Hits," but its value is "." indicating that it has been set to zero. The same is true for the "HmRun", "Runs", "RBI", "CAtBat", "CRBI", "CWalks", "DivisionW", "Errors", "Assists", and "NewLeagueN" variables. On the other hand, the following variables have non-zero coefficients, and thus, they are included in the final model: "Years", "CHits", "CHmRun", "CRuns", "CWalks", "LeagueA", "LeagueN", "PutOuts", "Hits:HmRun", "Hits:Walks", "Hits:CHits", "Hits:CRuns", "Hits:CRBI", "Hits:LeagueN", "Hits:DivisionW", "HmRun:LeagueN", "HmRun:Assists", "HmRun:NewLeagueN", "Runs:CHmRun", "RBI:CRuns", "Walks:CRBI". The lasso regression algorithm suggests that these variables are the most important for predicting the response variable in this model. However, we should be careful not to interpret these coefficients causally.